

Hardware Simulations of Open Quantum Systems by First-Order Hamiltonian Dilation

September 9, 2026

Abstract

This note gives a workflow for simulating Lindblad dynamics on IBM-style superconducting hardware using the first-order Hamiltonian dilation method of Ding–Li–Lin, together with mid-circuit measurement/reset and optional step-size/zero-noise extrapolation. The emphasis is not asymptotic complexity, but rather model problems that can be implemented and benchmarked on current devices: a two-qubit damped exchange model, a boundary-driven spin chain with measurable nonequilibrium current, and a quantum-detailed-balance thermalization model. The central practical feature is that one physical Lindblad step of size Δt is implemented by a coherent unitary whose dissipative coupling angles scale as $\sqrt{\Delta t}$, followed by tracing out or resetting the ancilla.

Contents

1	Conventions and target hardware workflow	2
2	Lindblad equation and first-order Hamiltonian dilation	2
2.1	The important $\sqrt{\Delta t}$ scale	3
2.2	Density-operator initialization	3
2.3	Generic first-order Trotterization	4
3	Model problem I: two-qubit damped exchange model	4
3.1	System model	4
3.2	Ancilla encoding and explicit dilated Hamiltonian	5
3.3	Trotter simplifications for the two-qubit model	6
3.4	Initial states and expected behavior	6
3.5	Measurements	6
3.6	Circuit Construction Possibilities	7
4	Model problem II: boundary-driven XXZ chain and steady current	7
4.1	Hamiltonian and boundary jumps	7
4.2	Dilated Hamiltonian	7
4.3	Trotter strategy	8
4.4	Initial states, evolution time, and steady-state certification	8
4.5	What to measure and how to interpret it	9
4.6	Relation to linear response and autocorrelations	9

5	Model problem III: quantum-detailed-balance thermalization	10
5.1	Target Gibbs state and QDB condition	10
5.2	Explicit QDB jump operators	10
5.3	Dilated Hamiltonian for the thermalization generator	11
5.4	Trotter simplifications for QDB jumps	11
5.5	Certification of thermal convergence	12
6	Benchmarking and error-mitigation plan	12
6.1	What has already been benchmarked	12
6.2	Three error sources to separate	12
6.3	Step-size extrapolation and ZNE	13
7	Concrete summer milestones	13
8	Implementation checklist	13

1 Conventions and target hardware workflow

We use the computational convention

$$Z|0\rangle = +|0\rangle, \quad Z|1\rangle = -|1\rangle, \quad (1)$$

so that

$$\sigma^- = |0\rangle\langle 1| = \frac{X + iY}{2}, \quad \sigma^+ = |1\rangle\langle 0| = \frac{X - iY}{2}. \quad (2)$$

Thus σ^- is the amplitude-damping/lowering operator from $|1\rangle$ to $|0\rangle$.

The hardware loop for one trajectory-free Lindblad simulation is:

Step 1. prepare the system state ρ_0 ;

Step 2. prepare the ancilla register in $|0\rangle_a$;

Step 3. apply a dilation unitary $U_{\Delta t} = \exp\left(-i\sqrt{\Delta t}\tilde{H}_{\Delta t}\right)$;

Step 4. measure and reset the ancilla, or equivalently trace it out;

Step 5. repeat for $N = T/\Delta t$ steps;

Step 6. measure system observables in appropriate Pauli bases.

No postselection is used: all ancilla outcomes are retained. The ancilla measurement record can be stored as a diagnostic of effective jump statistics, but the Lindblad estimator is the unconditional average over all outcomes.

2 Lindblad equation and first-order Hamiltonian dilation

The Markovian open-system dynamics is

$$\frac{d\rho}{dt} = \mathcal{L}(\rho) = -i[H, \rho] + \sum_{j=1}^J \left(V_j \rho V_j^\dagger - \frac{1}{2} \{V_j^\dagger V_j, \rho\} \right), \quad (3)$$

where $H = H^\dagger$ is the system Hamiltonian and V_j are jump operators. The first-order dilation used here follows the Ding–Li–Lin construction [6] and is closely related in spirit to earlier quantum algorithms for Lindblad simulation [4, 13].

For one time step Δt , define an ancilla space with basis states

$$|0\rangle_a, |1\rangle_a, \dots, |J\rangle_a. \quad (4)$$

The first-order dilated Hamiltonian is

$$\tilde{H}_{\Delta t} = \begin{bmatrix} \sqrt{\Delta t} H & V_1^\dagger & V_2^\dagger & \cdots & V_J^\dagger \\ V_1 & 0 & 0 & \cdots & 0 \\ V_2 & 0 & 0 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ V_J & 0 & 0 & \cdots & 0 \end{bmatrix}. \quad (5)$$

Equivalently,

$$\tilde{H}_{\Delta t} = \sqrt{\Delta t} |0\rangle\langle 0|_a \otimes H + \sum_{j=1}^J \left(|j\rangle\langle 0|_a \otimes V_j + |0\rangle\langle j|_a \otimes V_j^\dagger \right). \quad (6)$$

The corresponding channel is

$$\Phi_{\Delta t}(\rho) = \text{Tr}_a \left[U_{\Delta t} (|0\rangle\langle 0|_a \otimes \rho) U_{\Delta t}^\dagger \right], \quad U_{\Delta t} = \exp \left(-i\sqrt{\Delta t} \tilde{H}_{\Delta t} \right). \quad (7)$$

It satisfies

$$\Phi_{\Delta t}(\rho) = \rho + \Delta t \mathcal{L}(\rho) + O(\Delta t^2), \quad (8)$$

so $N = T/\Delta t$ repetitions give a first-order approximation over time T .

2.1 The important $\sqrt{\Delta t}$ scale

The exponent can be written as

$$\sqrt{\Delta t} \tilde{H}_{\Delta t} = \underbrace{\Delta t |0\rangle\langle 0|_a \otimes H}_{\text{coherent system term}} + \underbrace{\sqrt{\Delta t} \sum_{j=1}^J (|j\rangle\langle 0|_a \otimes V_j + \text{h.c.})}_{\text{dissipative dilation terms}}. \quad (9)$$

This is the central practical point. The physical Lindblad time step is Δt , but the jump-coupling rotations have angles proportional to $\sqrt{\Delta t}$. Therefore reducing Δt has two competing effects:

- smaller algorithmic error, $O(T\Delta t)$ for the first-order method;
- more repetitions $N = T/\Delta t$, hence more reset/readout/gate noise;
- only a $\sqrt{\Delta t}$ reduction in the dissipative rotation angles.

This makes the problem ideal for step-size extrapolation and hardware zero-noise extrapolation [14, 20, 11].

2.2 Density-operator initialization

On hardware, the density operator is not loaded as a matrix. We use one of the following procedures.

Pure initial state. Prepare $\rho_0 = |\psi_0\rangle\langle\psi_0|$ by an ordinary state-preparation circuit. This is the default for the two-qubit damping and current-buildup benchmarks.

Classical mixture. If

$$\rho_0 = \sum_{\ell} p_{\ell} |\psi_{\ell}\rangle\langle\psi_{\ell}|, \quad (10)$$

run separate jobs for the pure states $|\psi_{\ell}\rangle$ and combine measured expectations classically:

$$\langle O \rangle_{\rho_0}(t) = \sum_{\ell} p_{\ell} \langle O \rangle_{\psi_{\ell}}(t). \quad (11)$$

For diagonal Gibbs or infinite-temperature initial states, this is particularly efficient because $|\psi_{\ell}\rangle$ are computational-basis states.

Purification. A purification of ρ_0 can be used, but it costs extra qubits. For current IBM hardware, use classical mixture sampling unless a coherent reference system is essential.

2.3 Generic first-order Trotterization

Let

$$K_{\Delta t} := \sqrt{\Delta t} \tilde{H}_{\Delta t} = K_H + \sum_{j=1}^J K_j, \quad (12)$$

where

$$K_H = \Delta t |0\rangle\langle 0|_a \otimes H, \quad K_j = \sqrt{\Delta t} \left(|j\rangle\langle 0|_a \otimes V_j + |0\rangle\langle j|_a \otimes V_j^{\dagger} \right). \quad (13)$$

A first-order product formula is

$$e^{-iK_{\Delta t}} = e^{-iK_H} \prod_{j=1}^J e^{-iK_j} + O(\|K\|^2), \quad (14)$$

while a symmetric formula is

$$e^{-iK_{\Delta t}} = e^{-iK_H/2} \left(\prod_{j=1}^J e^{-iK_j/2} \right) \left(\prod_{j=J}^1 e^{-iK_j/2} \right) e^{-iK_H/2} + O(\|K\|^3). \quad (15)$$

Because the Lindblad method itself is first order, eq. (15) is mainly useful to reduce circuit-synthesis error so that the observed time-step error is dominated by the intended Lindblad discretization, not by the Hamiltonian-simulation subroutine.

3 Model problem I: two-qubit damped exchange model

3.1 System model

The first hardware target should be a two-qubit exchange Hamiltonian with local damping:

$$H_2 = \frac{J}{2}(X_1 X_2 + Y_1 Y_2) + \frac{\hbar}{2}(Z_1 + Z_2), \quad (16)$$

with jump operators

$$V_1 = \sqrt{\gamma} \sigma_1^-, \quad V_2 = \sqrt{\gamma} \sigma_2^-. \quad (17)$$

This is a minimal model with coherent excitation exchange and irreversible loss. It is small enough for exact classical reference simulation, and it exposes the essential hardware difficulty: repeated ancilla reset after dilation.

Useful starting parameters are

$$J = 1, \quad h = 0, \quad \gamma \in \{0.1, 0.2, 0.5\}, \quad \Delta t \in \{0.05, 0.1, 0.2\}. \quad (18)$$

Use dimensionless units with $\hbar = 1$. The expected relaxation time is $O(1/\gamma)$, so use

$$T \approx 2/\gamma \text{ to } 5/\gamma \quad (19)$$

for steady decay tests, but use shorter times $T \lesssim 1/\gamma$ for high-fidelity transient-current tests.

3.2 Ancilla encoding and explicit dilated Hamiltonian

There are $J = 2$ jumps, so the minimal ancilla has dimension 3. On qubit hardware, encode

$$|0\rangle_a \equiv |00\rangle_a, \quad |1\rangle_a \equiv |01\rangle_a, \quad |2\rangle_a \equiv |10\rangle_a, \quad (20)$$

leaving $|11\rangle_a$ unused. In the active qutrit subspace,

$$\tilde{H}_{\Delta t}^{(2)} = \begin{bmatrix} \sqrt{\Delta t} H_2 & V_1^\dagger & V_2^\dagger \\ V_1 & 0 & 0 \\ V_2 & 0 & 0 \end{bmatrix}. \quad (21)$$

Equivalently, defining $E_{mn} = |m\rangle \langle n|_a$,

$$K_{\Delta t}^{(2)} := \sqrt{\Delta t} \tilde{H}_{\Delta t}^{(2)} = \Delta t E_{00} \otimes H_2 + \sqrt{\Delta t} \sum_{q=1}^2 \left(E_{q0} \otimes V_q + E_{0q} \otimes V_q^\dagger \right). \quad (22)$$

For amplitude damping in our convention,

$$V_q = \frac{\sqrt{\gamma}}{2} (X_q + iY_q), \quad (23)$$

so if

$$A_{q0}^x = E_{q0} + E_{0q}, \quad A_{q0}^y = -i(E_{q0} - E_{0q}), \quad (24)$$

then

$$E_{q0} \otimes V_q + E_{0q} \otimes V_q^\dagger = \frac{\sqrt{\gamma}}{2} \left(A_{q0}^x \otimes X_q - A_{q0}^y \otimes Y_q \right). \quad (25)$$

The two terms in eq. (25) commute within the active two-level subspace $\text{span}\{|0\rangle_a, |q\rangle_a\}$, so each local damping block is an exchange-like $XX + YY$ rotation between the ancilla transition and the system qubit. This is a useful circuit optimization.

3.3 Trotter simplifications for the two-qubit model

For eq. (16), the terms X_1X_2 , Y_1Y_2 , and $Z_1 + Z_2$ mutually commute when the field is uniform. Therefore e^{-iK_H} can be compiled exactly as controlled Pauli rotations under the control projector E_{00} :

$$e^{-i\Delta t E_{00} \otimes H_2} = e^{-i\frac{J\Delta t}{2} E_{00} \otimes X_1 X_2} e^{-i\frac{J\Delta t}{2} E_{00} \otimes Y_1 Y_2} e^{-i\frac{h\Delta t}{2} E_{00} \otimes Z_1} e^{-i\frac{h\Delta t}{2} E_{00} \otimes Z_2}. \quad (26)$$

The jump terms K_1 and K_2 do not generally commute because they share the same ancilla state $|0\rangle_a$. Use the symmetric ordering

$$U_{\Delta t}^{\text{trotter}} = e^{-iK_H/2} e^{-iK_1/2} e^{-iK_2} e^{-iK_1/2} e^{-iK_H/2} \quad (27)$$

for the first hardware benchmark. If the compiled depth is too high, compare against the cheaper ordering $e^{-iK_H} e^{-iK_1} e^{-iK_2}$ and use exact classical simulation to separate Trotter error from device noise.

3.4 Initial states and expected behavior

Use the following initial states:

$$\rho_0^{(1)} = |10\rangle\langle 10|, \quad \text{localized excitation,} \quad (28)$$

$$\rho_0^{(2)} = |01\rangle\langle 01|, \quad \text{opposite localized excitation,} \quad (29)$$

$$\rho_0^{(+)} = |\psi_+\rangle\langle \psi_+|, \quad |\psi_+\rangle = \frac{|10\rangle + |01\rangle}{\sqrt{2}}, \quad \text{one-excitation eigenstate for } h = 0. \quad (30)$$

For $\gamma > 0$, all initial states should relax toward $|00\rangle\langle 00|$ in this convention. For $|10\rangle$ or $|01\rangle$, one should see coherent exchange oscillations damped by the loss. For $|\psi_+\rangle$, the exchange part is stationary when $h = 0$, so the signal is mostly damping.

3.5 Measurements

Measure:

$$n_q(t) = \frac{1 - \langle Z_q(t) \rangle}{2}, \quad q = 1, 2, \quad (31)$$

$$C_{XY}(t) = \langle X_1 X_2 + Y_1 Y_2 \rangle, \quad (32)$$

$$j_{12}(t) = J \langle X_1 Y_2 - Y_1 X_2 \rangle. \quad (33)$$

Here n_q is the excitation population on qubit q . The current j_{12} follows from the continuity equation for the exchange Hamiltonian in eq. (16). On hardware, j_{12} requires two Pauli-basis measurements: $X_1 Y_2$ and $Y_1 X_2$.

Expected signatures:

- $n_1 + n_2$ decays approximately exponentially on a time scale $1/\gamma$;
- $n_1 - n_2$ oscillates with frequency set by J and decays due to damping;
- $j_{12}(t)$ is transient for localized initial states and tends to zero at long times;
- step-size extrapolation in Δt should improve agreement with exact Lindblad evolution until hardware noise dominates.

3.6 Circuit Construction Possibilities

The following circuit constructions are possible for Model I, more details later : i) Unitary Circuits : Simplified for both symmetric and first order but quantum control operations due to projectors could lead to large depth ii) Dynamic Circuits : We need one quantum control and 2 classical controls for symmetric case but for first order we could just reduce it to one classical control iii) Qubit Reuse : For the first order, it could further be simplified by just using 1 ancilla rather than j ancillas using dynamic circuit approach.

4 Model problem II: boundary-driven XXZ chain and steady current

4.1 Hamiltonian and boundary jumps

The ambitious but still realistic target is a boundary-driven XXZ chain on $n = 2$ or $n = 3$ system qubits:

$$H_{\text{XXZ}} = \sum_{\ell=1}^{n-1} \left[\frac{J}{2} (X_{\ell} X_{\ell+1} + Y_{\ell} Y_{\ell+1}) + \frac{\Delta}{4} Z_{\ell} Z_{\ell+1} \right]. \quad (34)$$

The Z -magnetization current across bond $\ell, \ell + 1$ is

$$j_{\ell} = J (X_{\ell} Y_{\ell+1} - Y_{\ell} X_{\ell+1}), \quad (35)$$

with the above normalization.

The boundary reservoirs are represented by four jump operators:

$$V_{L,+} = \sqrt{\frac{\Gamma}{2}(1+\mu)} \sigma_1^+, \quad V_{L,-} = \sqrt{\frac{\Gamma}{2}(1-\mu)} \sigma_1^-, \quad (36)$$

$$V_{R,+} = \sqrt{\frac{\Gamma}{2}(1-\mu)} \sigma_n^+, \quad V_{R,-} = \sqrt{\frac{\Gamma}{2}(1+\mu)} \sigma_n^-. \quad (37)$$

For an isolated boundary spin, the left bath favors magnetization $+\mu$ while the right bath favors magnetization $-\mu$. Thus μ is the driving bias and Γ is the system-bath coupling. This model is standard in studies of nonequilibrium steady states of boundary-driven spin chains [15, 16, 8, 12].

Recommended initial parameter choices:

$$n = 2, 3, \quad J = 1, \quad \Delta \in \{0, 0.5, 1.5\}, \quad \Gamma \in \{0.1, 0.2, 0.5\}, \quad \mu \in \{0.1, 0.5\}. \quad (38)$$

Use $\mu = 0.1$ for linear-response comparisons and $\mu = 0.5$ for a larger hardware signal.

4.2 Dilated Hamiltonian

There are $J_{\text{jump}} = 4$ jumps, so the ancilla dimension is 5. Use a 3-qubit ancilla register and encode active states

$$|0\rangle_a, |L+\rangle_a, |L-\rangle_a, |R+\rangle_a, |R-\rangle_a, \quad (39)$$

leaving the other computational states unused. The first-order dilation is

$$\tilde{H}_{\Delta t}^{\text{bdry}} = \sqrt{\Delta t} |0\rangle\langle 0|_a \otimes H_{\text{XXZ}} + \sum_{\alpha \in \{L+, L-, R+, R-\}} \left(|\alpha\rangle\langle 0|_a \otimes V_{\alpha} + |0\rangle\langle \alpha|_a \otimes V_{\alpha}^{\dagger} \right). \quad (40)$$

The unitary per Lindblad step is

$$U_{\Delta t}^{\text{bdry}} = \exp\left(-i\sqrt{\Delta t} \tilde{H}_{\Delta t}^{\text{bdry}}\right). \quad (41)$$

For compilation, define

$$K_{\Delta t}^{\text{bdry}} = \sqrt{\Delta t} \tilde{H}_{\Delta t}^{\text{bdry}} = \Delta t |0\rangle\langle 0|_a \otimes H_{\text{XXZ}} + \sqrt{\Delta t} \sum_{\alpha} (|\alpha\rangle\langle 0|_a \otimes V_{\alpha} + \text{h.c.}). \quad (42)$$

Again, the Hamiltonian part has $O(\Delta t)$ angles, while each boundary jump coupling has $O(\sqrt{\Delta t \Gamma})$ angles.

4.3 Trotter strategy

For $n = 2$, all Hamiltonian terms in eq. (34) act on the same bond and the three Pauli strings XX , YY , and ZZ commute, so the controlled Hamiltonian block can be synthesized exactly as controlled Pauli rotations.

For $n \geq 3$, use an even-odd bond decomposition:

$$H_{\text{XXZ}} = H_{\text{even}} + H_{\text{odd}}, \quad (43)$$

where each part is a sum of disjoint two-qubit bond Hamiltonians. A symmetric step is

$$e^{-iK_{\Delta t}^{\text{bdry}}} \approx e^{-iK_H/2} \left(\prod_{\alpha} e^{-iK_{\alpha}/2} \right) \left(\prod_{\alpha}^{\text{reverse}} e^{-iK_{\alpha}/2} \right) e^{-iK_H/2}, \quad (44)$$

where $K_H = \Delta t |0\rangle\langle 0|_a \otimes H_{\text{XXZ}}$ and K_{α} is the α th boundary jump block. The order of the four boundary jumps should be fixed across all time steps; reversing the order in the second half gives a Strang-type cancellation of leading synthesis error.

Possible simplifications:

- for $n = 2$, compile the full active-subspace unitary exactly and compare with Pauli product formulas;
- for $\Delta = 0$, the XX model has fewer noncommuting Hamiltonian terms;
- for weak drive $\mu \ll 1$, collect data for both $+\mu$ and $-\mu$ and antisymmetrize the current to cancel even-in- μ biases;
- use Pauli grouping for $X_{\ell}X_{\ell+1}$ and $Y_{\ell}Y_{\ell+1}$ rotations, which commute on each bond.

4.4 Initial states, evolution time, and steady-state certification

Good initial states are:

$$\rho_0 = |00 \cdots 0\rangle\langle 00 \cdots 0|, \quad (45)$$

$$\rho_0 = 2^{-n} \mathbb{I}, \quad \text{implemented by uniform computational-basis sampling,} \quad (46)$$

$$\rho_0 = |10 \cdots 0\rangle\langle 10 \cdots 0|, \quad \text{for current-buildup tests.} \quad (47)$$

The steady-state time should be estimated classically for each n and parameter set by diagonalizing the small Liouvillian or by exact time integration. Let λ_{gap} denote the Liouvillian spectral gap, i.e. the smallest nonzero decay rate. A practical target is

$$T \gtrsim \frac{3}{\lambda_{\text{gap}}} \quad \text{for a visible plateau,} \quad T \gtrsim \frac{5}{\lambda_{\text{gap}}} \quad \text{for quantitative steady-state estimates.} \quad (48)$$

For hardware, also run shorter transients because the current buildup may be more robust than the final long-time plateau.

Certify stationarity by checking:

$$|\langle Z_j(t + \Delta t) \rangle - \langle Z_j(t) \rangle| \ll 1, \quad (49)$$

$$|\langle j_\ell(t + \Delta t) \rangle - \langle j_\ell(t) \rangle| \ll 1, \quad (50)$$

$$\langle j_1 \rangle \approx \langle j_2 \rangle \approx \dots \approx \langle j_{n-1} \rangle \quad \text{in the bulk for } n \geq 3. \quad (51)$$

4.5 What to measure and how to interpret it

Measure the magnetization profile and current:

$$m_j = \langle Z_j \rangle, \quad \langle j_\ell \rangle = J (\langle X_\ell Y_{\ell+1} \rangle - \langle Y_\ell X_{\ell+1} \rangle). \quad (52)$$

Expected behavior:

- $m_1 > m_n$ for $\mu > 0$;
- the steady current changes sign when $\mu \mapsto -\mu$;
- for small μ , $\langle j \rangle$ should be approximately linear in μ ;
- the current magnitude depends on Δ , with integrable/ballistic-like behavior near the XX regime and more suppressed transport for stronger anisotropy in small systems.

4.6 Relation to linear response and autocorrelations

In the weak-driving regime $\mu \ll 1$ and weak system–bath coupling $\Gamma \ll J$, the boundary-driven nonequilibrium steady current is a finite-size open-system proxy for a linear-response transport coefficient. In closed equilibrium systems, the Green–Kubo relation expresses transport coefficients in terms of equilibrium current autocorrelations,

$$\kappa_L(\beta) \sim \int_0^\infty C_{JJ}(t) dt, \quad C_{JJ}(t) = \frac{1}{L} \langle J_{\text{tot}}(t) J_{\text{tot}}(0) \rangle_\beta, \quad (53)$$

where

$$J_{\text{tot}} = \sum_{\ell=1}^{n-1} j_\ell. \quad (54)$$

The equality between a boundary-driven finite-size NESS current and the closed-system Green–Kubo coefficient is not automatic. It is a comparison to make in the linear-response, weak-coupling, finite-size regime. Recent studies of boundary-driven spin chains explicitly discuss this connection and its limitations [8, 12]. For the summer project, the practical goal is to measure the NESS current and compare its small- μ slope against exact classical Lindblad results and, where feasible, against a closed-system autocorrelation calculation.

5 Model problem III: quantum-detailed-balance thermalization

5.1 Target Gibbs state and QDB condition

For thermalization, choose a model whose target Gibbs state is known:

$$\rho_\beta = \frac{e^{-\beta H_{\text{th}}}}{Z_\beta}, \quad Z_\beta = \text{Tr}(e^{-\beta H_{\text{th}}}). \quad (55)$$

A Lindblad generator satisfies quantum detailed balance (QDB) with respect to ρ_β if it is self-adjoint under the appropriate ρ_β -weighted inner product, such as the GNS/KMS inner product. A simple hardware-friendly construction is a diagonal, reversible, single-spin-flip dynamics. This is a classical detailed-balance Markov chain embedded into a Lindblad equation, and it is the right first experiment before attempting noncommuting Gibbs samplers.

Take the two-qubit Ising Hamiltonian

$$H_{\text{th}} = -h(Z_1 + Z_2) - K Z_1 Z_2. \quad (56)$$

For $z_1, z_2 \in \{+1, -1\}$, define

$$E(z_1, z_2) = -h(z_1 + z_2) - K z_1 z_2. \quad (57)$$

The target Gibbs probabilities are

$$\pi_\beta(z_1, z_2) = \frac{e^{-\beta E(z_1, z_2)}}{\sum_{z_1, z_2} e^{-\beta E(z_1, z_2)}}. \quad (58)$$

5.2 Explicit QDB jump operators

Let

$$P_s^{(q)} = \frac{\mathbb{I} + s Z_q}{2}, \quad s \in \{+1, -1\}. \quad (59)$$

Conditioned on the neighboring spin value s , flipping qubit q from $|1\rangle$ to $|0\rangle$ or from $|0\rangle$ to $|1\rangle$ changes the Ising energy by a known amount. With our convention $Z|0\rangle = +|0\rangle$ and $Z|1\rangle = -|1\rangle$, define bounded detailed-balance rates

$$r_s^+ = \eta \frac{e^{\beta(h+Ks)}}{2 \cosh[\beta(h+Ks)]}, \quad r_s^- = \eta \frac{e^{-\beta(h+Ks)}}{2 \cosh[\beta(h+Ks)]}. \quad (60)$$

Then

$$\frac{r_s^+}{r_s^-} = e^{2\beta(h+Ks)}, \quad (61)$$

which is exactly the Boltzmann ratio for the two directions of the conditional spin flip.

Use eight conditional jump operators:

$$V_{1,s}^+ = \sqrt{r_s^+} \sigma_1^+ P_s^{(2)}, \quad V_{1,s}^- = \sqrt{r_s^-} \sigma_1^- P_s^{(2)}, \quad (62)$$

$$V_{2,s}^+ = \sqrt{r_s^+} P_s^{(1)} \sigma_2^+, \quad V_{2,s}^- = \sqrt{r_s^-} P_s^{(1)} \sigma_2^-, \quad s = \pm 1. \quad (63)$$

For diagonal density matrices, this is exactly a reversible continuous-time Markov chain with stationary distribution π_β . Since each forward/backward transition pair satisfies eq. (61), the generator

obeys detailed balance. The corresponding embedded Lindblad dynamics has ρ_β as a stationary state. If coherent off-diagonal components are present in ρ_0 , add optional dephasing jumps

$$V_{\phi,q} = \sqrt{\gamma_\phi} Z_q, \quad (64)$$

which commute with H_{th} and leave ρ_β invariant.

Recommended parameters:

$$h = 0.3, \quad K = 0.7, \quad \beta \in \{0.5, 1, 2\}, \quad \eta \in \{0.1, 0.2, 0.5\}. \quad (65)$$

The case $\beta = 0$ should be included as a sanity check: the target is the maximally mixed state and all conditional forward/backward rates are equal.

5.3 Dilated Hamiltonian for the thermalization generator

Ignoring optional dephasing, there are $J_{\text{jump}} = 8$ jumps. The active ancilla dimension is 9, so use four ancilla qubits. Label the eight active nonzero ancilla states by

$$\alpha \in \{(1, +, +), (1, +, -), (1, -, +), (1, -, -), (2, +, +), (2, +, -), (2, -, +), (2, -, -)\}, \quad (66)$$

where the entries denote site, flip direction, and neighbor condition. Let V_α denote the corresponding operator in eqs. (62) and (63). The first-order dilated Hamiltonian is

$$\tilde{H}_{\Delta t}^{\text{th}} = \sqrt{\Delta t} |0\rangle\langle 0|_a \otimes H_{\text{th}} + \sum_{\alpha=1}^8 \left(|\alpha\rangle\langle 0|_a \otimes V_\alpha + |0\rangle\langle \alpha|_a \otimes V_\alpha^\dagger \right). \quad (67)$$

The unitary step is

$$U_{\Delta t}^{\text{th}} = \exp\left(-i\sqrt{\Delta t} \tilde{H}_{\Delta t}^{\text{th}}\right). \quad (68)$$

If dephasing jumps are included, append them to the same sum, increasing the active ancilla dimension accordingly.

5.4 Trotter simplifications for QDB jumps

The Hamiltonian H_{th} is diagonal, so $e^{-i\Delta t|0\rangle\langle 0|_a \otimes H_{\text{th}}}$ consists only of controlled Z_1 , Z_2 , and $Z_1 Z_2$ rotations. Each conditional jump is a projected amplitude-raising or amplitude-lowering exchange. For example,

$$|\alpha\rangle\langle 0|_a \otimes \sqrt{r_s^+} \sigma_1^+ P_s^{(2)} + \text{h.c.} \quad (69)$$

acts only when the neighbor qubit is in the $Z = s$ sector. Thus it can be compiled as a controlled exchange between the ancilla transition $|0\rangle_a \leftrightarrow |\alpha\rangle_a$ and the system qubit, with an additional control on $P_s^{(2)}$.

Useful optimizations:

- jumps with orthogonal projectors $P_+^{(q)}$ and $P_-^{(q)}$ can be scheduled in either order because they act on disjoint system subspaces;
- the two Pauli strings inside each σ^+ or σ^- exchange commute as in the amplitude-damping example;
- if hardware resources are tight, first run a one-qubit Gibbs sampler with two jumps $\sqrt{r^+}\sigma^+$ and $\sqrt{r^-}\sigma^-$, then the two-qubit conditional sampler.

5.5 Certification of thermal convergence

Because H_{th} is diagonal, certification can avoid full tomography. Measure computational-basis populations

$$p_t(z_1, z_2) = \Pr_t(Z_1 = z_1, Z_2 = z_2) \quad (70)$$

and compare against $\pi_\beta(z_1, z_2)$. Recommended metrics are

$$d_{\text{TV}}(t) = \frac{1}{2} \sum_{z_1, z_2} |p_t(z_1, z_2) - \pi_\beta(z_1, z_2)|, \quad (71)$$

$$D_{\text{KL}}(p_t \| \pi_\beta) = \sum_{z_1, z_2} p_t(z_1, z_2) \log \frac{p_t(z_1, z_2)}{\pi_\beta(z_1, z_2)}, \quad (72)$$

$$E(t) = \sum_{z_1, z_2} p_t(z_1, z_2) E(z_1, z_2), \quad (73)$$

$$C_{12}(t) = \langle Z_1 Z_2 \rangle_t. \quad (74)$$

For a two-qubit experiment, also perform occasional full state tomography to check that coherences are small. A practical convergence fit is

$$|E(t) - E_\beta| \approx Ae^{-\lambda t}, \quad (75)$$

where λ is an empirical mixing rate. Choose $T \gtrsim 3/\lambda$ for a first plateau and $T \gtrsim 5/\lambda$ for quantitative comparison.

The thermalization experiment connects to recent work on dissipative Gibbs-state preparation and quantum detailed balance [17, 2, 1, 18]. The proposed two-qubit model is deliberately simpler than those many-body Gibbs samplers, but it tests exactly the hardware primitive needed by those algorithms: repeated, unconditional simulation of a detailed-balanced Lindbladian.

6 Benchmarking and error-mitigation plan

6.1 What has already been benchmarked

Related hardware demonstrations include QITE-based simulation of Lindblad dynamics on IBM hardware [10], adaptive variational Lindblad simulation on IBM quantum processors [3], and mid-circuit/dynamic-circuit benchmarking and qubit reuse [7, 19, 9, 5]. These works motivate the present project but do not replace it: the specific benchmark here is the first-order Hamiltonian dilation with repeated ancilla reset, explicit $\sqrt{\Delta t}$ dissipative rotations, and physical observables such as steady current and Gibbs convergence.

6.2 Three error sources to separate

For each model, report at least three comparisons:

1. exact Lindblad integration versus exact first-order dilation channel;
2. exact first-order dilation channel versus Trotterized/synthesized dilation unitary on a noiseless simulator;
3. noisy hardware result versus noiseless Trotterized circuit.

This separation is crucial. Otherwise one cannot tell whether failure comes from Lindblad time discretization, Hamiltonian-simulation/Trotter error, reset/readout error, or ordinary gate noise.

6.3 Step-size extrapolation and ZNE

For a scalar observable O , collect estimates

$$f(\Delta t) = \text{Tr}[O\Phi_{\Delta t}^{T/\Delta t}(\rho_0)]. \quad (76)$$

For the first-order method,

$$f(\Delta t) = f(0) + c_1\Delta t + c_2\Delta t^2 + \dots, \quad (77)$$

until hardware noise dominates. Use Richardson or least-squares extrapolation over several Δt values. Chebyshev-like step grids should be tested because they can reduce interpolation variance compared with equally spaced grids [14].

At each Δt , also consider zero-noise extrapolation by gate folding or calibrated noise scaling. The desired two-stage estimator is

$$f_{\text{ideal}}(0) \approx \lim_{\Delta t \rightarrow 0} \lim_{\lambda_{\text{noise}} \rightarrow 0} f_{\text{hardware}}(\Delta t, \lambda_{\text{noise}}). \quad (78)$$

A practical question is whether the order of these extrapolations matters within statistical uncertainty.

7 Concrete summer milestones

Weeks 1–2: software reference. Implement exact Lindblad integration, exact dilation channel, and Trotterized dilation circuits for the two-qubit model. Produce plots of $n_1(t)$, $n_2(t)$, $C_{XY}(t)$, and $j_{12}(t)$.

Weeks 3–4: first IBM runs. Run the two-qubit damping model with both fresh ancillas and mid-circuit reset. Quantify reset/readout overhead by comparing them at the same effective T and Δt .

Weeks 5–7: step-size extrapolation and ZNE. Run $\Delta t \in \{0.2, 0.1, 0.05\}$ or a Chebyshev-like grid. Compare raw, step-extrapolated, ZNE-only, and combined estimators.

Weeks 8–10: boundary-driven current. Run the $n = 2$ boundary-driven model first, then $n = 3$ if depth allows. Target the sign, magnitude, and μ -linearity of $\langle j \rangle$.

Weeks 11–12: QDB thermalization. Run the one-qubit Gibbs sampler as a calibration, then the two-qubit conditional-jump sampler. Certify convergence by total-variation distance to the exact Gibbs distribution and by energy relaxation.

8 Implementation checklist

- Fix the Pauli and raising/lowering convention before coding.
- Build $\tilde{H}_{\Delta t}$ symbolically and verify that $\Phi_{\Delta t}(\rho) = \rho + \Delta t\mathcal{L}(\rho) + O(\Delta t^2)$ numerically.
- Validate every compiled $U_{\Delta t}$ against exact matrix exponentiation for one step.

- Always record ancilla outcomes, but do not postselect on them.
- Use the same transpilation and layout strategy across all Δt values in an extrapolation study.
- For every plotted hardware observable, include exact Lindblad, exact dilation, noiseless Trotter, and hardware data when feasible.
- Report shot counts and uncertainty bars; extrapolation without error bars is not meaningful.

References

- [1] Eric Brunner et al. Lindblad engineering for quantum gibbs state preparation under the eigenstate thermalization hypothesis. *Quantum*, 9:1843, 2025.
- [2] Chi-Fang Chen, Michael J. Kastoryano, Fernando G. S. L. Brandão, and András Gilyén. Quantum thermal state preparation. *arXiv preprint arXiv:2303.18224*, 2023.
- [3] Huo Chen, Niladri Gomes, Siyuan Niu, and Wibe Albert de Jong. Adaptive variational simulation for open quantum systems. *Quantum*, 8:1252, 2024.
- [4] Richard Cleve and Chunhao Wang. Efficient quantum algorithms for simulating lindblad evolution. In *44th International Colloquium on Automata, Languages, and Programming (ICALP 2017)*, volume 80 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 17:1–17:14, Dagstuhl, Germany, 2017. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
- [5] Matthew DeCross, Eli Chertkov, Megan Kohagen, and Michael Foss-Feig. Qubit-reuse compilation with mid-circuit measurement and reset. *Physical Review X*, 13:041057, 2023.
- [6] Zhiyan Ding, Xiantao Li, and Lin Lin. Simulating open quantum systems using hamiltonian simulations. *PRX Quantum*, 5(2):020332, 2024.
- [7] L. C. G. Govia et al. A randomized benchmarking suite for mid-circuit measurements. *arXiv preprint arXiv:2207.04836*, 2024.
- [8] T. Heitmann et al. Spin-1/2 XXZ chain coupled to two lindblad baths: Constructing nonequilibrium steady states from equilibrium correlation functions. *Physical Review B*, 108:L201119, 2023.
- [9] Fei Hua, Yuwei Jin, Yanhao Chen, Suhas Vittal, Kevin Krsulich, Lev S. Bishop, John Lapeyre, Ali Javadi-Abhari, and Eddy Z. Zhang. Exploiting qubit reuse through mid-circuit measurement and reset. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, 2023. Also available as arXiv:2211.01925.
- [10] Hirsh Kamakari, Shi-Ning Sun, Mario Motta, and Austin J. Minnich. Digital quantum simulation of open quantum systems using quantum imaginary-time evolution. *PRX Quantum*, 3:010320, 2022.
- [11] Abhinav Kandala, Kristan Temme, Antonio D. Córcoles, Antonio Mezzacapo, Jerry M. Chow, and Jay M. Gambetta. Error mitigation extends the computational reach of a noisy quantum processor. *Nature*, 567:491–495, 2019.

- [12] M. Kraft et al. Lindblad dynamics from spatio-temporal correlation functions in boundary-driven quantum spin chains. *Physical Review Research*, 6:023251, 2024.
- [13] Xiantao Li and Chunhao Wang. Simulating markovian open quantum systems using higher-order series expansion. In *50th International Colloquium on Automata, Languages, and Programming (ICALP 2023)*, volume 261 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 87:1–87:20, Dagstuhl, Germany, 2023. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
- [14] Pegah Mohammadipour and Xiantao Li. Reducing circuit depth in lindblad simulation via step-size extrapolation. *arXiv preprint arXiv:2507.22341*, 2025.
- [15] Tomaž Prosen. Open XXZ spin chain: Nonequilibrium steady state and a strict bound on ballistic transport. *Physical Review Letters*, 106:217206, 2011.
- [16] Tomaž Prosen and Berislav Buča. Integrable non-equilibrium steady state density operators for boundary driven XXZ spin chains: Observables and full counting statistics. *Journal of Physics A: Mathematical and Theoretical*, 48:395003, 2015.
- [17] Patrick Rall, Chunhao Wang, and Pawel Wocjan. Thermal state preparation via rounding promises. *Quantum*, 7:1132, 2023.
- [18] Matteo Scandi et al. Thermalization in open many-body systems and KMS detailed balance. *arXiv preprint arXiv:2505.20064*, 2025.
- [19] L. Shirizly et al. Randomized benchmarking protocol for dynamic circuits. *Physical Review A*, 111:012611, 2025.
- [20] Kristan Temme, Sergey Bravyi, and Jay M. Gambetta. Error mitigation for short-depth quantum circuits. *Physical Review Letters*, 119:180509, 2017.